

How Autocracy Shapes the Social Sciences

Synthesized Evidence from an Agent Orchestra

Tore Wig

University of Oslo · AUTOKNOW (ERC-COG 101231399)

Workshop draft, June 2026; preliminary, please do not cite



Funded by
the European Union



European Research Council
Established by the European Commission



Does autocracy bend the *content* of social science?

- Autocracies are now major science producers (China > US).
- Old claim: **open societies** breed good science (Popper, 1945; Kitcher, 2001).
- Bibliometrics shows regime effects on **volume** and **citations** (Tayebi and Teimouri, 2025; Yang and Liu, 2021).

The gap

If researchers self-censor, autocracy should also shape the **contents and framing** of research, not just its quantity. We look at this in Social Sciences and Humanities (SSH) field.

The project at a glance

Research questions

1. Do researchers in autocracies **self-censor** and respond strategically to regime context?
2. What are the implications for **social-science contents**?

Design

- 25 **pre-registered** AI-agent teams
- One corpus: **2.7M** WoS SSH articles (1970–2023)
- Autocracy-democracy = V-Dem liberal democracy
- 5 mechanism families, family-wise Bonferroni corrections

So far (21/25 teams)

- **Essentially no self-censorship**
- 0 of 21 significant
- Method: an **agent-orchestra** many-analysts design

Preliminary: 21 of 25 teams done; 4 framing/embedding teams pending.

Theory: strategic self-censorship

Researchers and their options.

- Driven by **career** (prominence, tenure, funding) and **scientific** goals (move the field, solve puzzles).
- But also value physical, economic, and social **security**, and act strategically within regime-constraints
- Rationally to regime **sticks** (sanction, lost funding) and **carrots** (state-controlled research money).
- Strategic margins: *what* to study, *how* to frame it, *who* to collaborate with, how *visible* to be.

No single decisive test.

- Regimes' "red lines" are uncertain, to rulers and researchers alike.
- Several substitutable levers to pull (framing, contents, collaboration, language, etc.)
- Self-censorship therefore surfaces as **many weak signals**, spread across domains and dependent on context.

Test many implications on many outcomes, not one.

Five families of observable implications

Topic avoidance	avoid sensitive subjects, keywords, fields.	7
Framing neutrality	hedge, depoliticize, go technocratic.	6
Collaboration constraint	fewer international ties.	5
Visibility suppression	censored work gets cited less.	4
Ideological alignment	regime-consonant content.	3

Each → a directional prediction on Liberal democracy.


Pre-registration

- One theory, ~ 25 operationalizations, one corpus.
- One analyst $\Rightarrow$ huge **researcher degrees of freedom**.
- Cherry-picking

Response

Pre-register **many** implementations, lock them before results, and read the whole **distribution** of estimates.


One theory, many implications, one team each



h of the 25 teams tests offspring implication; the theory is judged on the *distribution* of results.

The Agent Orchestra

25 teams, each a pipeline of AI agents, with a PI gate at every stage (human in the loop):



Amber: PI review gate (redesign / revise loop). Hypotheses locked before analysis at Git 960525f; an emerging method (Zhu and Wang, 2026; Gupta et al., 2026).

Data: one locked corpus

Corpus (Web of Science, SSH)

- 2.7M articles, 167 countries
- 1970–2023
- 99.98% matched to V-Dem

Politics (V-Dem)

- IV: v2x_libdem
- Controls: log GDPpc, log pop

Unit: **country–year**.

One shared benchmark specification

Two-way fixed effects, every team:

$$Y_{it} = \beta \text{v2x_libdem}_{it} + \gamma_1 \log \text{gdppc}_{it} + \gamma_2 \log \text{pop}_{it} + \alpha_i + \delta_t + \varepsilon_{it}$$

- Country + year FE; SE clustered by country.
- Secondary: pooled OLS. Robustness: alt. measure + binary regime.
- **Family-wise Bonferroni**, $\alpha = 0.05/k$.

Example 1, Team 04: topic avoidance

Rationale. Regime-sensitive concepts such as democracy, human rights, corruption, and protest are the highest-cost topics under autocracy. A paper's most visible labels (title, keywords) are the likeliest site of **anticipatory avoidance**.

- **Outcome** `share_regime_sensitive`: share of a country-year's articles with ≥ 1 hit from a regime-sensitive dictionary (democracy, rights, corruption, protest, repression, ...) in title, keywords, or abstract.
- **Test:** TWFE on `v2x_libdem`; **predicted** $\beta > 0$ (more democracy $\rightarrow$ more sensitive terms).

Result

$\hat{\beta} = +0.001$ (SE 0.018), $p = 0.94$, $N = 2,803$ country-years.

Indistinguishable from zero.

Example 2, Team 24: ideological alignment

Rationale. Autocracies do not only suppress dissent; **reward** scholarship that endorses the regime (state funding, party-controlled outlets, promotion criteria).


- **Outcome** share_legitimizing: share of abstracts an **LLM** (Claude Haiku) labels as actively endorsing the political system; stratified sample, ~141k abstracts classified.
- **Test:** TWFE on v2x_libdem; **predicted** $\beta < 0$ (more autocracy $\rightarrow$ more legitimation).

Result

$\hat{\beta} = +0.003$ (SE 0.002), $p = 0.23$, $N = 3,207$ country-years. **Wrong sign, not significant.**

Only 582 of 140,879 classified abstracts (0.4%) coded as legitimating: a rare-event outcome. Completed after the May-27 forest plot was frozen, not shown there.

Results: coefficients cluster at zero



Each row = one pre-registered team test (TWFE: country and year FE, GDP/pop controls, clustered by country).

Solid line: zero. Dashed: nominal ± 1.96 . Dotted: family-specific Bonferroni 5% threshold.

Sign convention: positive z = outcome rises with liberal democracy (typically the theory-predicted direction).

16 of 25 pre-registered teams reported here; 7 LLM/embedding-based teams (05, 06, 10, 12, 23–25) pending.

Preliminary: 21 of 25 teams done; 4 framing/embedding teams pending.

Nothing survives correction

Sub-family	Registered (k)	Completed tests	Sign of $\hat{\beta}$		Nominal $p < .05$	Bonferroni $p < .05$
			> 0	< 0		
Topic avoidance	7	5	4	1	0	0
Framing neutrality	6	4	3	1	0	0
Collaboration constraint	5	5	4	1	0	0
Visibility suppression	4	4	1	3	1	0
Total	22	18	12	6	1	0

Four corpus-based families (18 teams) shown.

- **0 of 18** survive Bonferroni; 1 nominal hit (T22) runs **opposite** to prediction.
- Ideological-alignment teams **23–25** (now complete) are **also null**: $p = 0.13, 0.23, 0.16$.
- Our two examples land here too: **T04** $p = 0.94$, **T24** $p = 0.23$.

Is the null robust, or just under-powered?

Robustness

- Within-family directions suggestive at best; none significant.
- TWFE and pooled OLS **agree**: not a fixed-effects artifact.

Power, the leading alternative explanation

- Large samples (2,500–3,500 country-years, 150+ countries): well-powered for **between-country** contrasts.
- But TWFE leans on **within-country** change in democracy, which is slow and rare, so the within estimand is less precisely identified.
- Reassurance: confidence intervals sit **tight around zero**, ruling out large effects, and the cross-sectional (pooled) estimates agree.
- **Still open**: small effects, or effects concentrated in specific regimes, fields, or periods, cannot be excluded.

What does a null mean?

- Across 21 dimensions, including LLM-coded **ideological alignment**, **little sign** that autocracy bends published content.
- **Caveats:**
 - Observational; content censorship vs. **venue selection**.
 - WoS skews English / high-impact; misses domestic outlets.
 - **4 framing teams** (LLM, embeddings) still pending.

Preliminary: 21 of 25 teams done; 4 framing/embedding teams pending.

Summary and discussion

Findings. Across 21 pre-registered tests, including LLM-coded ideological alignment, **little evidence** that autocracy bends the content of published SSH; no coefficient survives correction.

Contribution

- *Theoretical:* a first broad test of content self-censorship; an informative null that shifts the burden of proof.
- *Methodological:* an **agent-orchestra** many-analysts design that scales a multi-implication theory, with pre-registration and human review built in.

Limitations. Observational (content censorship vs. venue selection); WoS skews English and high-impact, missing domestic outlets; effective power for the within estimand is limited; 4 framing teams pending; AI-generated design choices raise provenance and refereeing questions.

► Appendix / backup slides

tore.wig@stv.uio.no

References I

- Gupta, Jishu Sen, Harini SI, Somesh Kumar Singh, Syed Mohamad Tawseeq, Yaman Kumar Singla, David Doermann, Rajiv Ratn Shah and Balaji Krishnamurthy. 2026. "Accelerating Social Science Research via Agentic Hypothesization and Experimentation." *arXiv preprint arXiv:2602.07983* .
- Kitcher, Philip. 2001. *Science, truth, and democracy*. Oxford University Press.
- Popper, Karl R. 1945. *The Open Society and Its Enemies*. London: Routledge.
- Tayebi, Amir and Sheida Teimouri. 2025. "Democratization and knowledge in social sciences." *Public Choice* 202(1):77–108.
- Yang, Yi and Lin Liu. 2021. "The Politics of Academic Innovation: A Cross-National Study of the Effects of Regime Type on Knowledge Production." *Asian Journal of Technology Innovation* 29(3).
- Zhu, Chen and Xiaolu Wang. 2026. "HLER: Human-in-the-Loop Economic Research via Multi-Agent Pipelines for Empirical Discovery." *arXiv preprint arXiv:2603.07444* .

Appendix: backup slides

- ▶ All 25 teams: outcomes and status
- ▶ Five families and Bonferroni thresholds
- ▶ Estimation: specification and robustness
- ▶ Data and corpus construction
- ▶ LLM classification (ideological alignment)
- ▶ Pre-registration, gates, pending teams

All 25 teams: outcomes and status

Topic avoidance

- 01 PCI (political-content index)
- 02 Keyword entropy
- 03 Share in sensitive fields
- 04 Share w/ regime-sensitive terms
- 05 Share critical of own governance[†]
- 06 Semantic diversity (embeddings)[†]
- 07 Share internationally funded

Framing neutrality

- 08 Epistemic hedging rate
- 09 Normative language rate
- 10 Share purely technocratic[†]
- 11 Argumentative stance rate
- 12 Share w/ normative conclusion[†]
- 13 Share in domestic journals

Collaboration constraint

- 14 Mean co-author countries
- 15 Share democratic co-authors
- 16 Share domestic-only authorship
- 17 Mean author count
- 18 Mean distinct institutions

Visibility suppression

- 19 Log citation ratio
- 20 Citation gap by sensitivity
- 21 Share ever cited
- 22 Citation Gini

Ideological alignment

- 23 Share liberal/neutral framing
- 24 Share legitimating regime
- 25 Share anti-liberal framing

[†] pending (LLM / embeddings); the other 21 are complete.

◀ back

Five families and Bonferroni thresholds

Sub-family	Teams	k	Adjusted α ($0.05/k$)
Topic avoidance	01–07	7	$p < 0.0071$
Framing neutrality	08–13	6	$p < 0.0083$
Collaboration constraint	14–18	5	$p < 0.0100$
Visibility suppression	19–22	4	$p < 0.0125$
Ideological alignment	23–25	3	$p < 0.0167$
Total	01–25	25	

Bonferroni applied within each family; predicted signs are in each team's preregistration.

[◀ back](#)

Estimation: specification and robustness

Primary (TWFE), every team:

$$Y_{it} = \beta \text{v2x_libdem}_{it} + \gamma_1 \log \text{gdppc}_{it} + \gamma_2 \log \text{pop}_{it} + \alpha_i + \delta_t + \varepsilon_{it}$$

- α_i country FE, δ_t year FE; SE clustered by country.
- **Secondary:** pooled OLS, year FE only (cross-sectional pattern).
- **RC1:** alternative operationalization or sample restriction.
- **RC2:** replace `v2x_libdem` with `lied_binary`.
- Unit: country-year; minimum 10 articles per country-year (most teams).

◀ back

Data and corpus construction

- Source: Web of Science, all SSH subject categories.
- 3,189,557 article-country rows; 2,709,224 unique articles.
- Years 1970–2023; 167 countries; 99.98% matched to V-Dem by ISO3.
- De-duplication and V-Dem merge: `scripts/00_prepare_data.R`.
- IV `v2x_libdem` (0–1); controls log GDP per capita, log population.
- Robustness IV `lied_binary`; regime type `v2x_regime` (0–3).
- Corpus locked before any analysis.

◀ back

LLM classification (ideological alignment)

Teams 23–25 code abstracts with an LLM (Claude Haiku). Example, Team 24:

- **Binary label.** LEGITIMATING: the abstract explicitly endorses or validates the political system, state authority, ruling party, or leadership as effective, beneficial, or necessary. NON-LEGITIMATING: everything else (default when in doubt).
- **Sample.** Stratified by country $\times$ regime $\times$ decade; up to 10 abstracts per stratum; abstracts of ≥ 50 characters; fixed seed.
- **Aggregation.** Share LEGITIMATING per country-year, then TWFE as usual.
- **Scale.** About 141k abstracts classified for Team 24.

[◀ back](#)

Pre-registration, gates, pending teams

Pre-registration. All 25 hypotheses, predicted signs, primary specifications, robustness checks, and Bonferroni families were locked at Git commit 960525f (6 May 2026), before any analyst session ran.

- **Gates.** PI review at design (B), results (D), and report (G); revision loops allowed for implementation only, not for outcomes.
- **Complete (21):** all corpus-based teams plus ideological alignment (23–25).
- **Pending (4):** 05 (critical governance, LLM), 06 (semantic diversity, embeddings), 10 (technocratic share, LLM), 12 (normative conclusion, LLM).

[← back](#)